Log-Monitoring mit graylog

Thomas Darimont
Java User Group Saarland - 28. Treffen

16.02.2017

Graylog Website

https://www.graylog.org

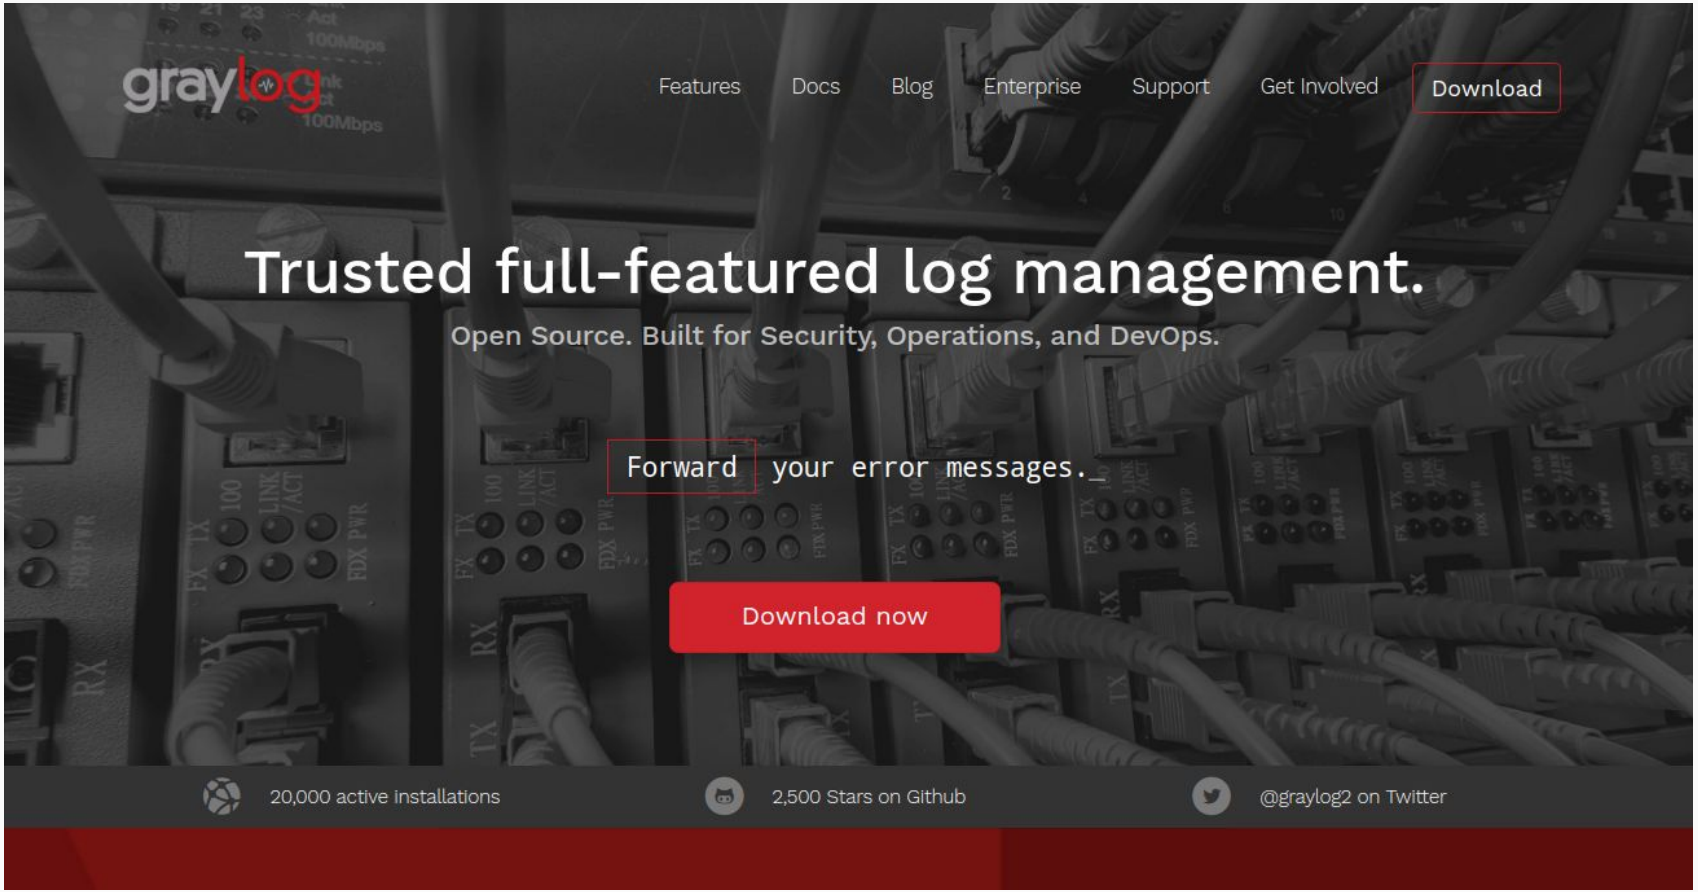

Graylog in a Nutshell

- Log Management Platform
- Collect, Index and Analyze Log data
- Structured and Unstructured
- Java based, Open Source GPLv3
- Uses Elasticsearch & MongoDB
- Multi-User

What is Graylog?

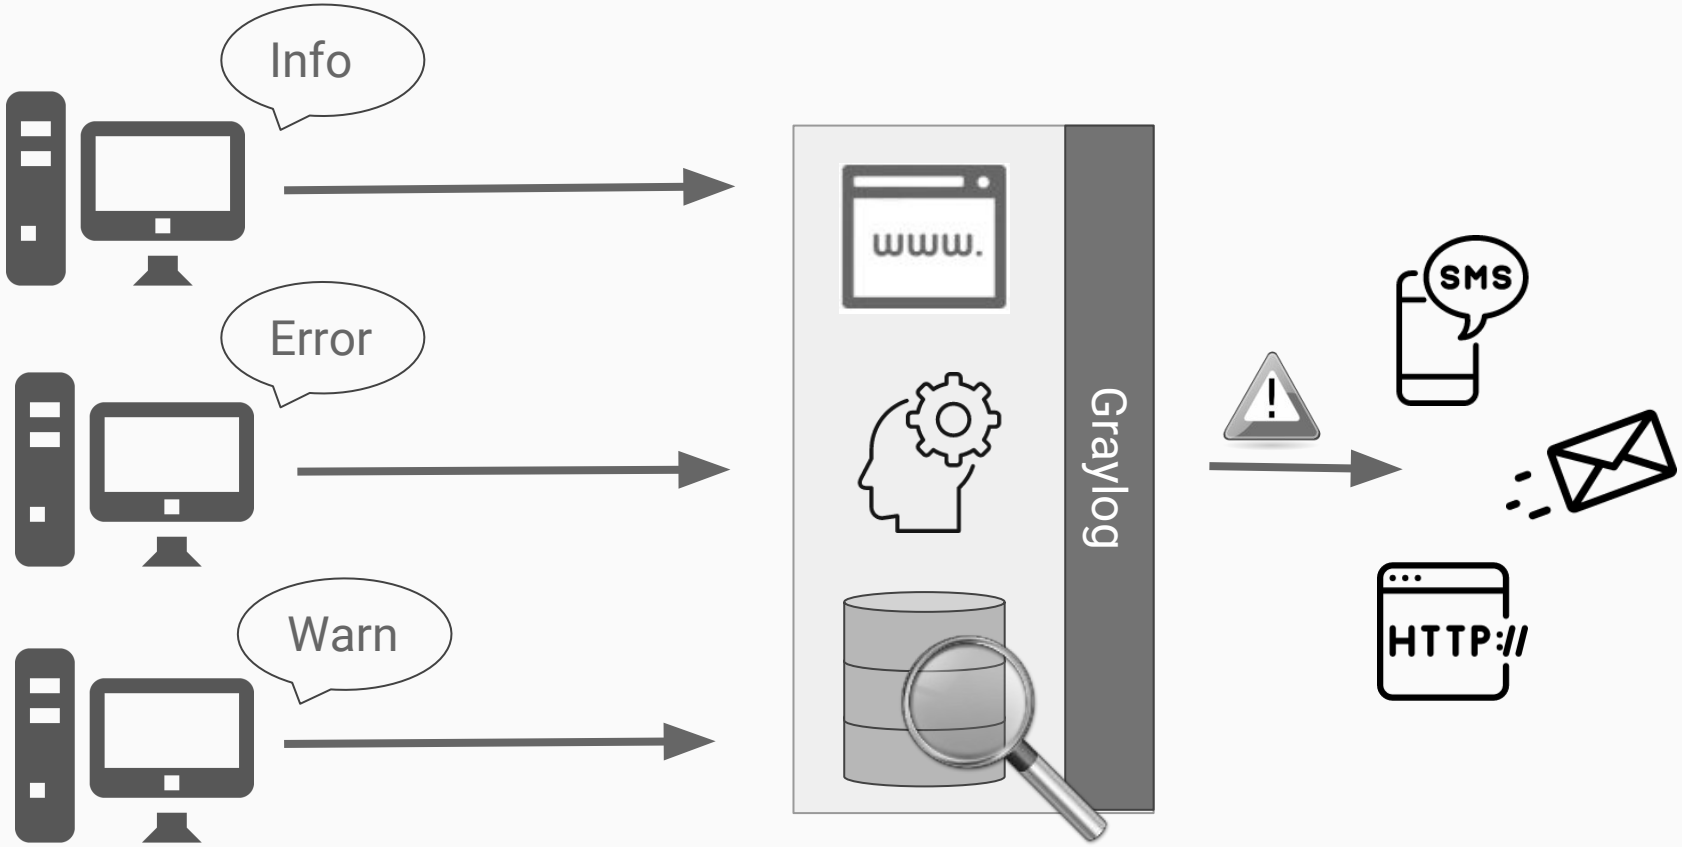

Graylog Facts

- Current Version 2.2.0 (Released 14. February)
- Very mature project > 6 years
- Docker, OVA Appliance, Standalone
- Free & Commercial (Graylog Inc.)
- Free Version quite powerful
- Enterprise: Support, Audit-Trail, Archiving++
- Trusted by Leading Companies (> 20.000 Installs)
- Graylog Marketplace

Graylog Marketplace

https://marketplace.graylog.org

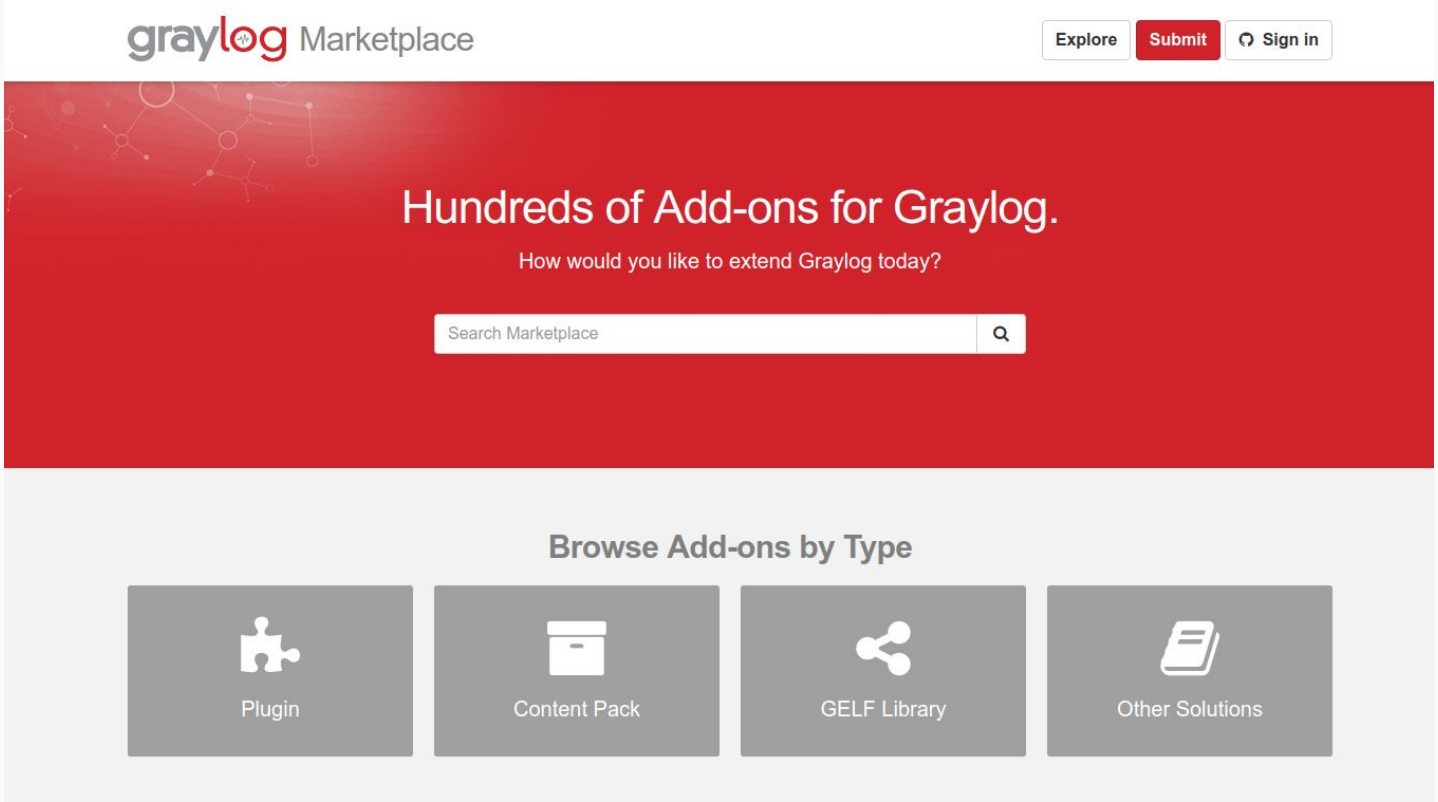

Graylog on Github

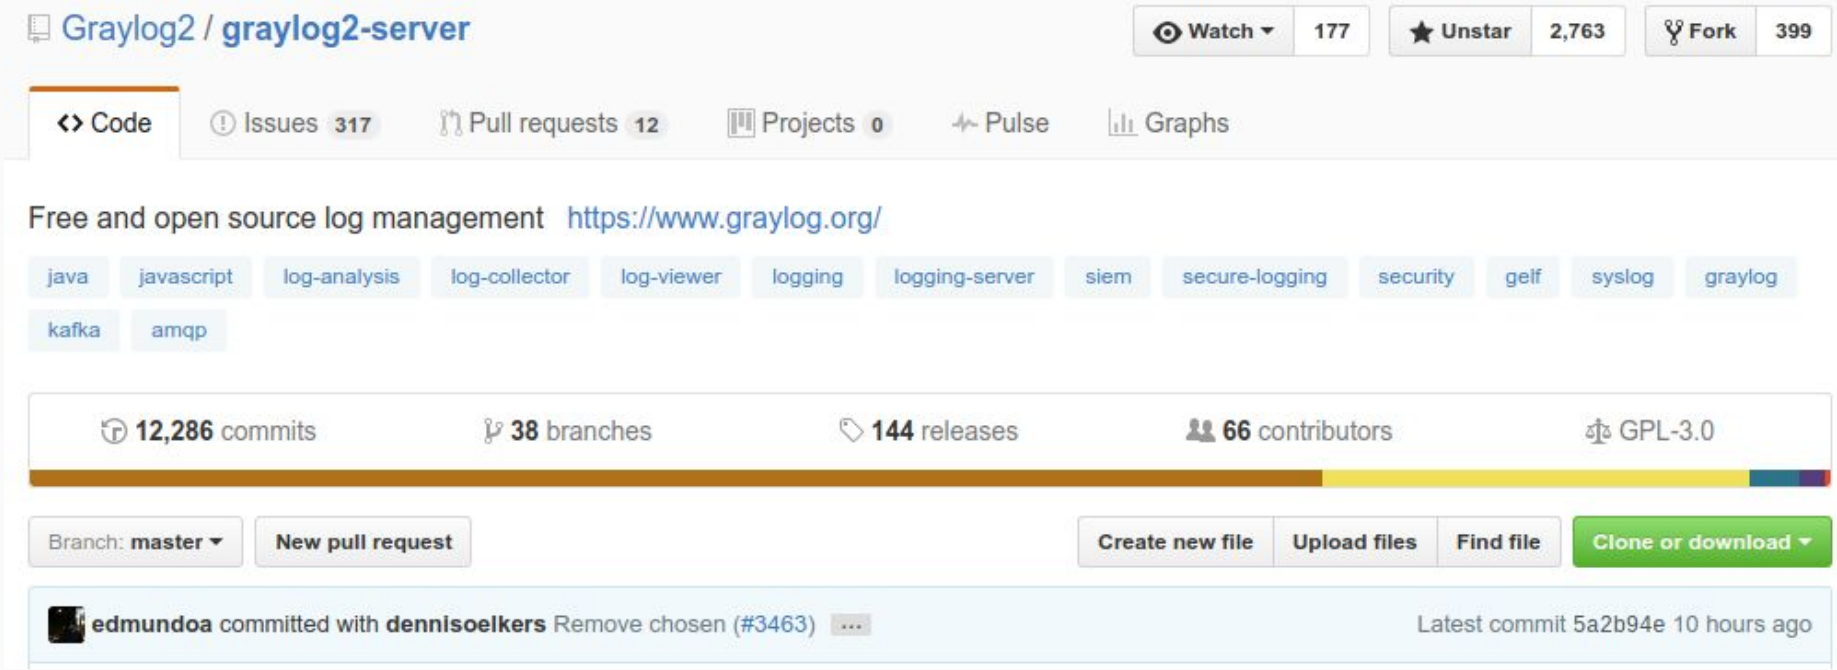

https://github.com/Graylog2/graylog2-server

Graylog Features

- Multiple Formats
  - SYSLOG, GELF, Beats, JSON, Plaintext, Raw,...
- Multiple Protocols
  - TCP, UDP, HTTP, AMQP, Kafka, ...
- Log Message Classification
- User Management and Access Control
- Scalable Architecture with HA support
- High Performance Log Processing

Graylog Concepts

- System — Config, Nodes, Indices, AuthN
- Inputs — Endpoints for receiving log data
- Indices — Store log data, controls log retention
- Streams — Rule based message routing & filtering
- Dashboards — Aggregated views on log data
- Alerts — Conditionally trigger & send notifications
- Outputs — Forward log data
- Pipelines — Stackable Pipes & Filters for log processing

Graylog Component Overview

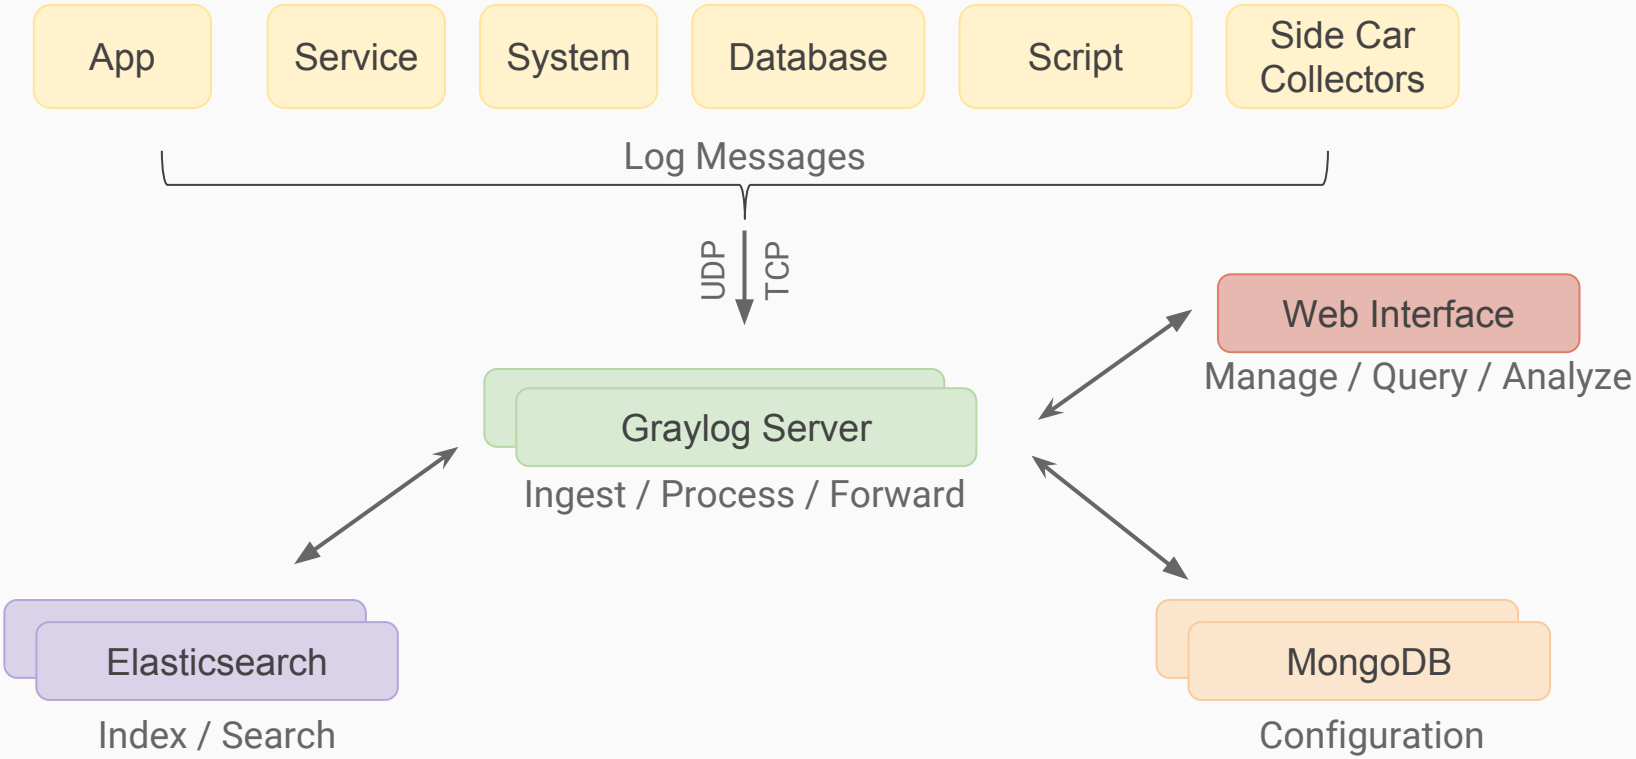

Graylog Canonical HA-Setup

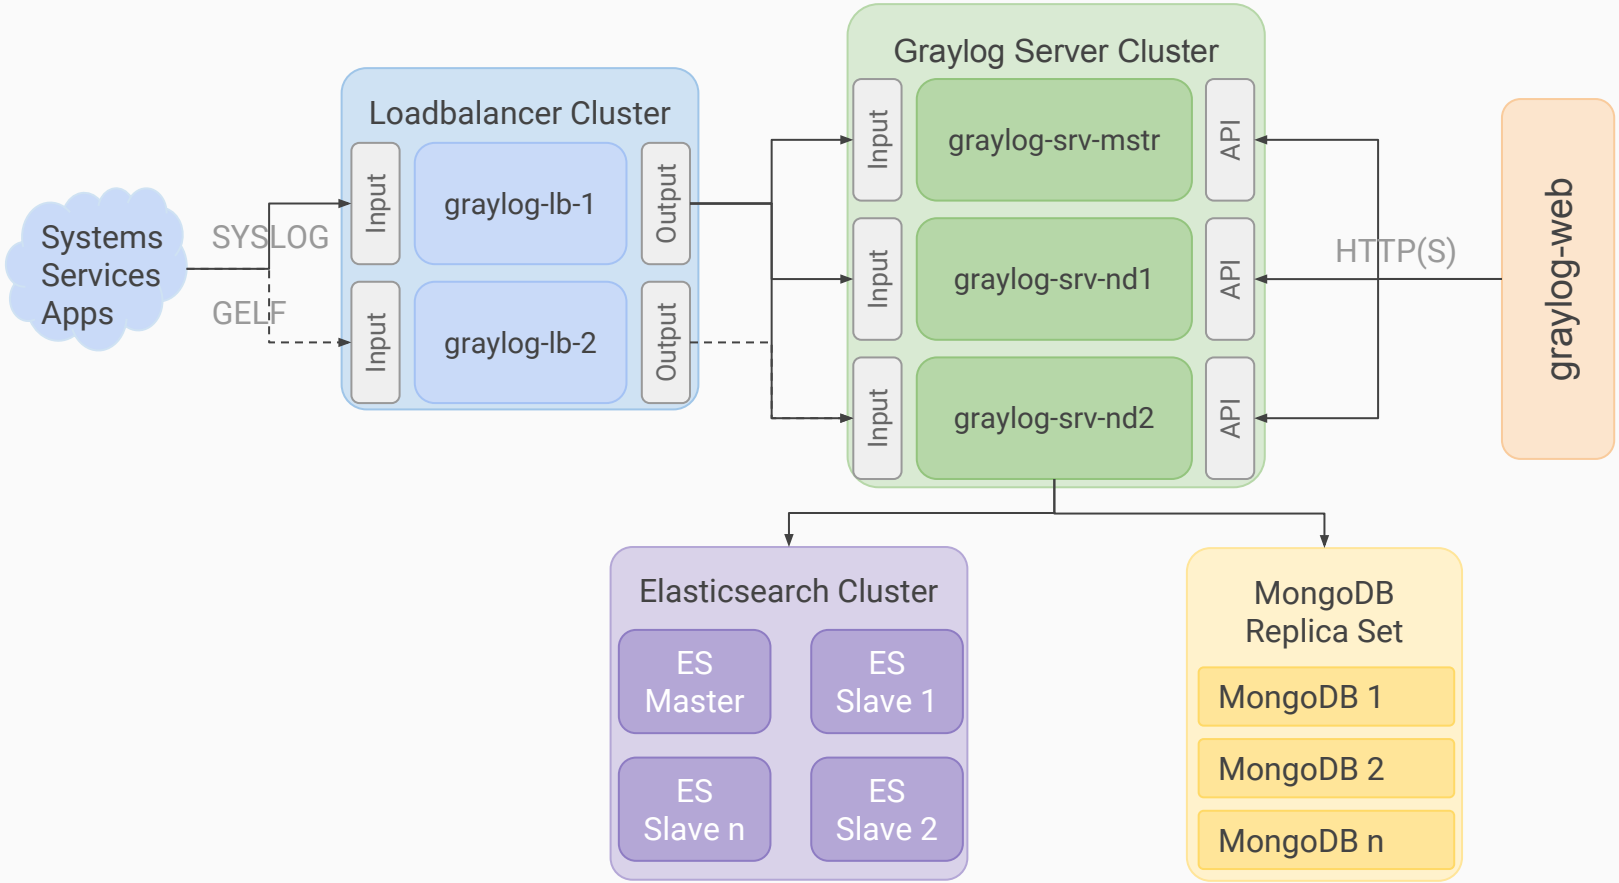

Graylog Canonical HA-Setup

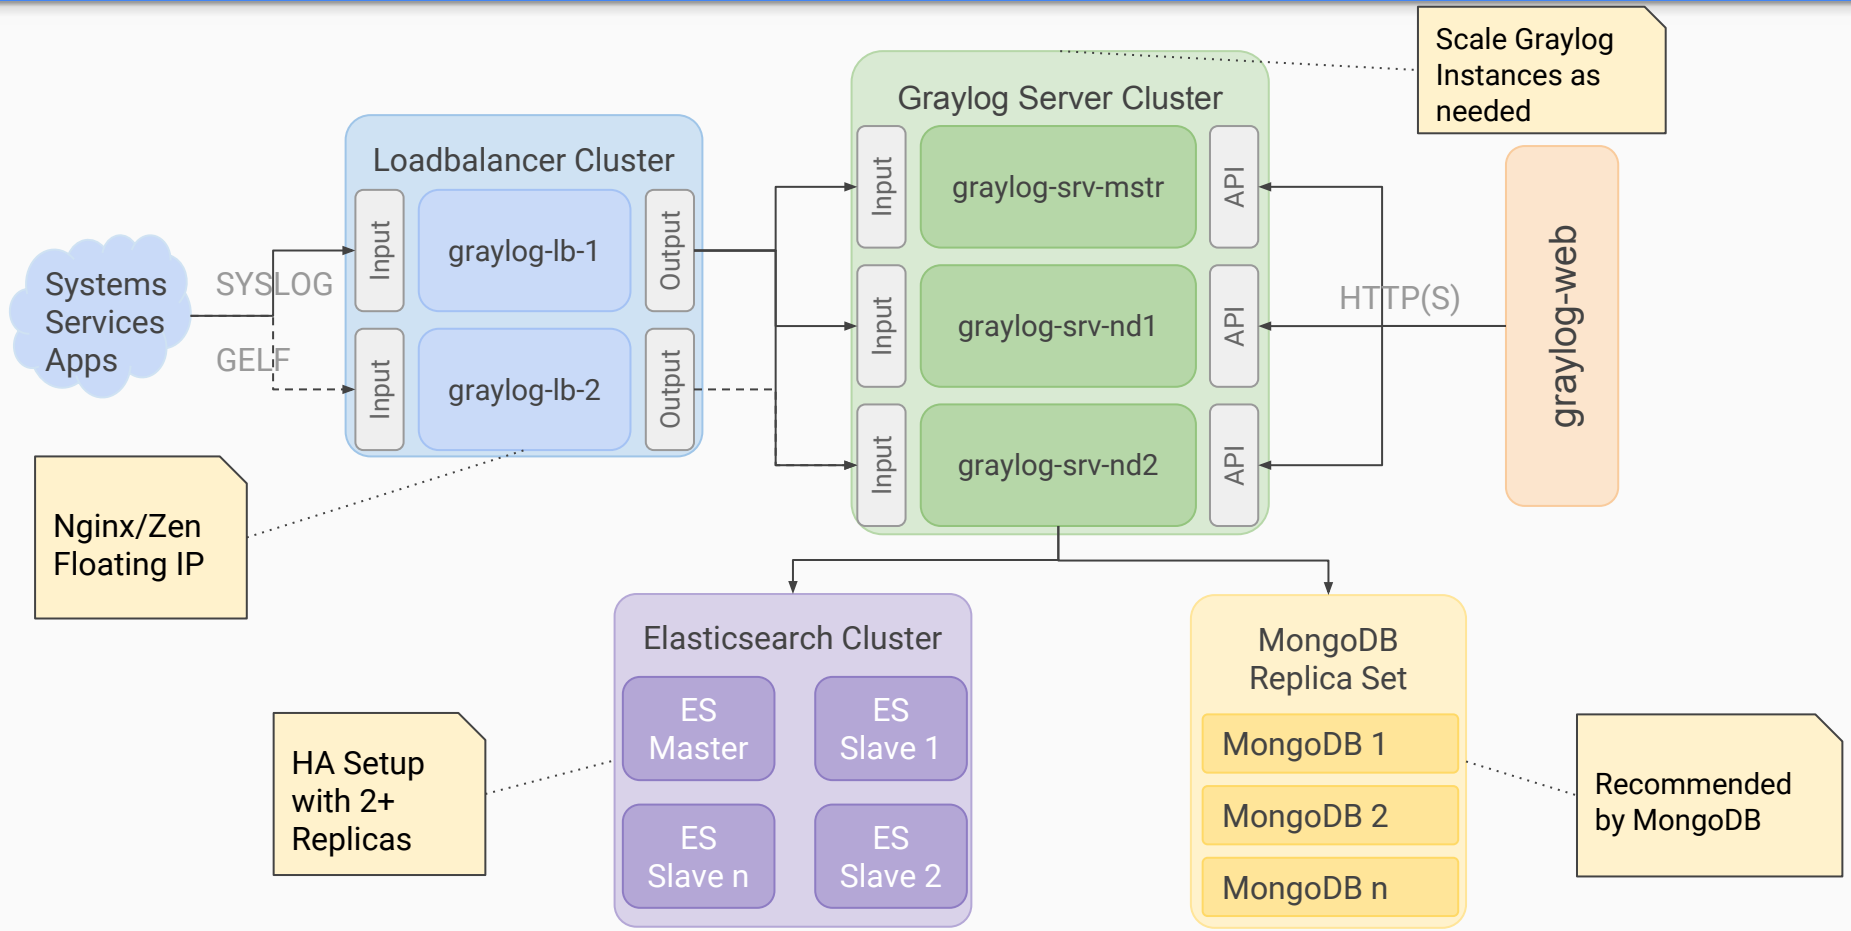

Graylog Architecture

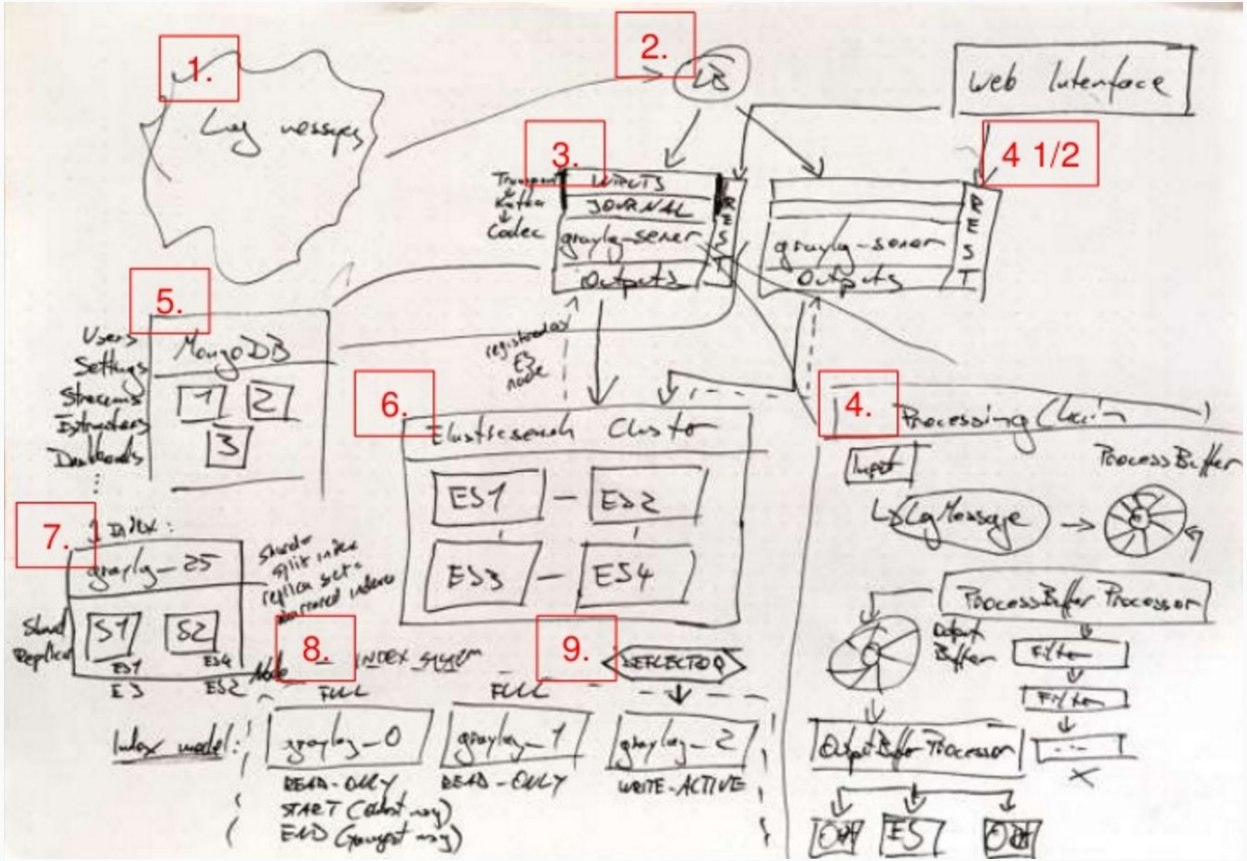

1. Log Messages
2. Load Balancer
3. Transport Layer
4. Processing Chain

4.1/2 REST API

5. MongoDB ReplicaSet
6. Elasticsearch Cluster
7. Anatomy of an Index
8. Index Model
9. Deflector Queue

[Graylog Engineering Design your Architecture]

Interesting Architecture Bits...

- Uses Apache Kafka for the append-only log journal on disk
  - Allows fast writes to disk
  - Avoids losing messages during spikes
- Uses LMAX Disruptor RingBuffer
  - Allows fast data ingestion and processing with low-latency
- Graylog Node acts as non-data Elasticsearch Node
  - Allows faster native protocols instead of HTTP/JSON
- Designed for Horizontal Scalability and HA
  - Graylog Nodes (2n+1 Processor nodes)
  - MongoDB (Shards + Replicas)
  - Elasticsearch (Shards + Replicas)
- Frontend build with React
- Custom Log Format GELF for more flexibility

Graylog Extended Log Format

- JSON String
- Avoids shortcomings of classic plain syslog
- Structured Log Message with Types
- Supports custom fields
- UDP and TCP
- Chunking
- Compression
- … GELF reference

```
{
  "version": "1.1",
  "host": "example.org",
  "short_message": "A short message",
  "full_message": "Backtrace here\n\nmore stuff",
  "timestamp": 1385053862.3072,
  "level": 1,
  "_user_id": 9001,
  "_some_info": "foo",
  "_some_env_var": "bar"
}
```

Demo Send GELF Message from a Shell Script

```
#!/usr/bin/env bash

script_execution_id=$(uuidgen)

log_gelf(){
    msg=$1
    nc -w 1 -u logserver.tdlabs.local 12205 <<EOF
{
  "version":"1.1"
, "host":"$(hostname)"
, "short_message":"$msg"
, "full_message":"$msg"
, "level":1
, "_script_execution_id":"$script_execution_id"
}\0
EOF
}

log_gelf "Hello from $0"
```

GELF with
netcat and heredoc
Gist

Demo use GELF logging with Docker

```
docker run -dit \
           --name nginx \
           -p 28080:80 \
             --log-driver=gelf \
             --log-opt gelf-address=udp://logserver.tdlabs.local:12205 \
             nginx:1.11.9-alpine
```

See: https://docs.docker.com/engine/admin/logging/overview

DEMO

Graylog in Action

Recap

- System
- Inputs
- Streams
- Searches
- Dashboards
- Alerts
- REST API Browser

Inputs

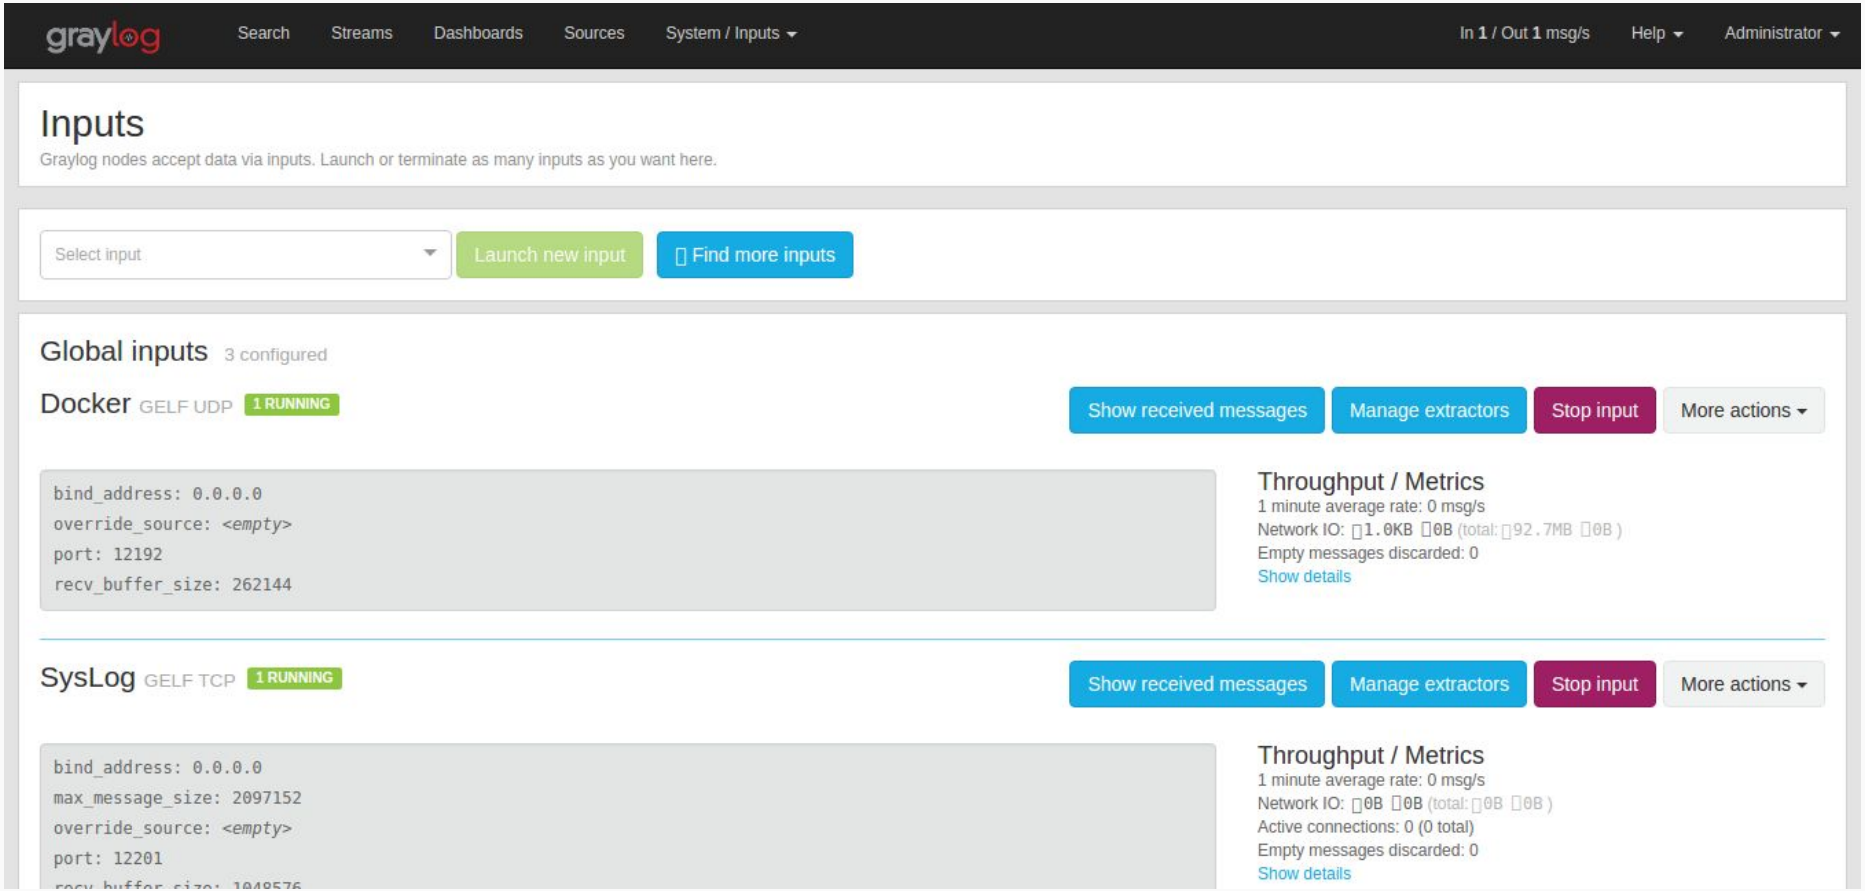
graylog
Search
Streams
Dashboards
Sources
System / Inputs
In 1 / Out 1 msg/s
Help
Administrator
Inputs
Graylog nodes accept data via inputs. Launch or terminate as many inputs as you want here.
Select input
Launch new input
Find more inputs
Global inputs 3 configured
Docker GELF UDP 1 RUNNING
Show received messages
Manage extractors
Stop input
More actions
bind_address: 0.0.0.0
override_source: <empty>
port: 12192
recv_buffer_size: 262144
Throughput / Metrics
1 minute average rate: 0 msg/s
Network IO: 1.0KB 0B (total: 92.7MB 0B)
Empty messages discarded: 0
Show details
SysLog GELF TCP 1 RUNNING
Show received messages
Manage extractors
Stop input
More actions
bind_address: 0.0.0.0
max_message_size: 2097152
override_source: <empty>
port: 12201
Throughput / Metrics
1 minute average rate: 0 msg/s
Network IO: 0B 0B (total: 0B 0B)
Active connections: 0 (0 total)
Empty messages discarded: 0
Show details

Streams

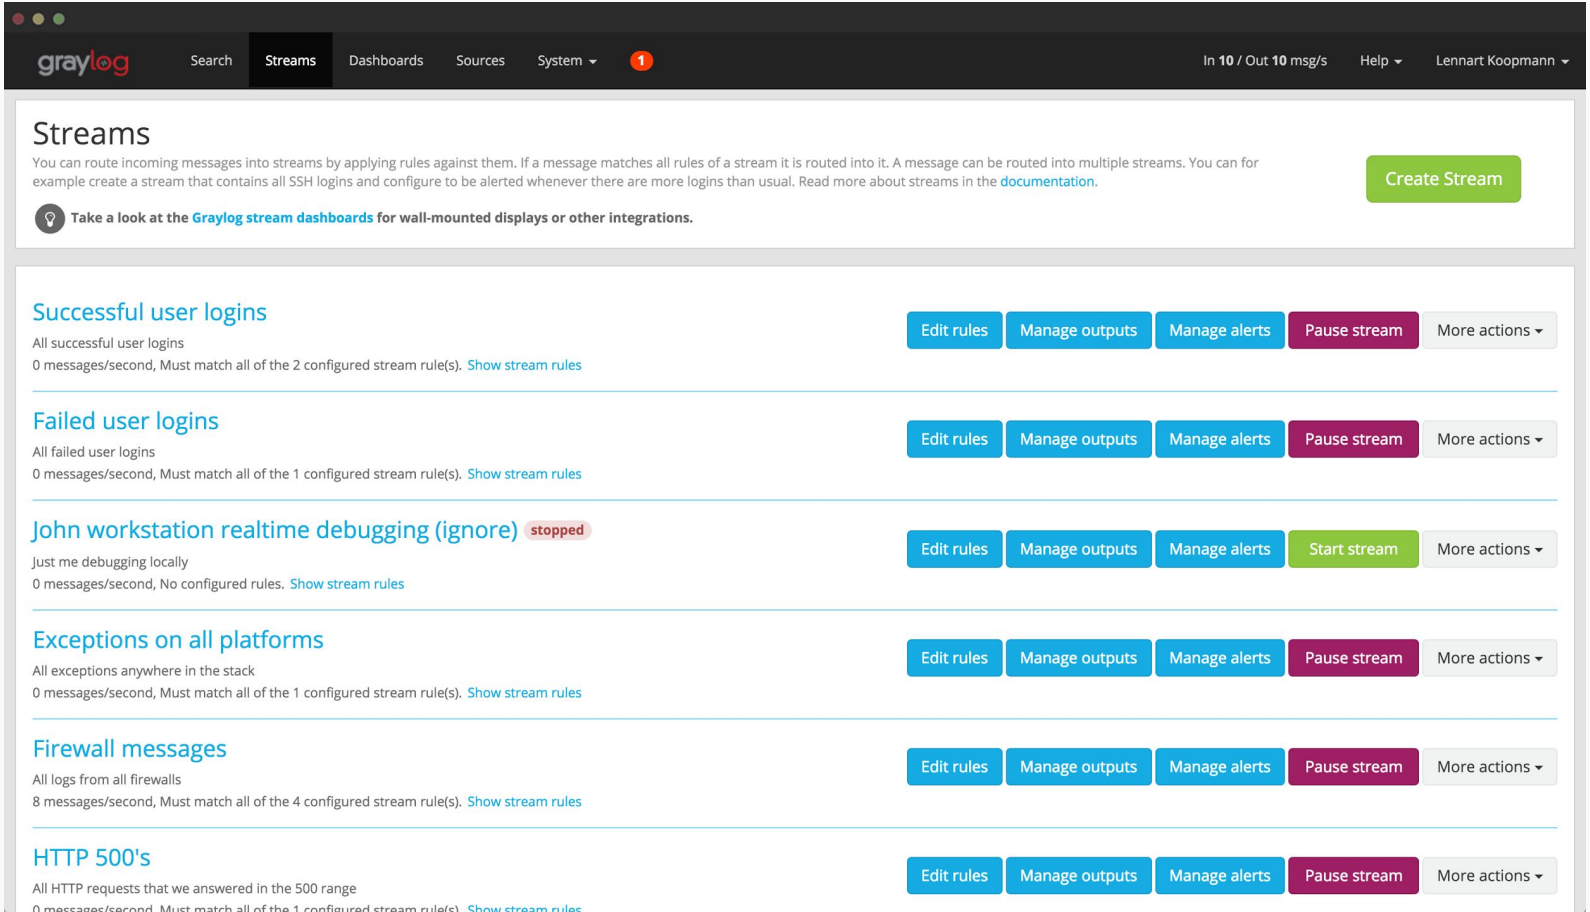
graylog
Search
Streams
Dashboards
Sources
System
1
In 10 / Out 10 msg/s
Help
Lennart Koopmann
Streams
You can route incoming messages into streams by applying rules against them. If a message matches all rules of a stream it is routed into it. A message can be routed into multiple streams. You can for example create a stream that contains all SSH logins and configure to be alerted whenever there are more logins than usual. Read more about streams in the documentation.
Create Stream
Take a look at the Graylog stream dashboards for wall-mounted displays or other integrations.
Successful user logins
All successful user logins
0 messages/second, Must match all of the 2 configured stream rule(s). Show stream rules
Edit rules
Manage outputs
Manage alerts
Pause stream
More actions
Failed user logins
All failed user logins
0 messages/second, Must match all of the 1 configured stream rule(s). Show stream rules
Edit rules
Manage outputs
Manage alerts
Pause stream
More actions
John workstation realtime debugging (ignore)
stopped
Just me debugging locally
0 messages/second, No configured rules. Show stream rules
Edit rules
Manage outputs
Manage alerts
Start stream
More actions
Exceptions on all platforms
All exceptions anywhere in the stack
0 messages/second, Must match all of the 1 configured stream rule(s). Show stream rules
Edit rules
Manage outputs
Manage alerts
Pause stream
More actions
Firewall messages
All logs from all firewalls
8 messages/second, Must match all of the 4 configured stream rule(s). Show stream rules
Edit rules
Manage outputs
Manage alerts
Pause stream
More actions
HTTP 500's
All HTTP requests that we answered in the 500 range
Edit rules
Manage outputs
Manage alerts
Pause stream
More actions

Log Message Search

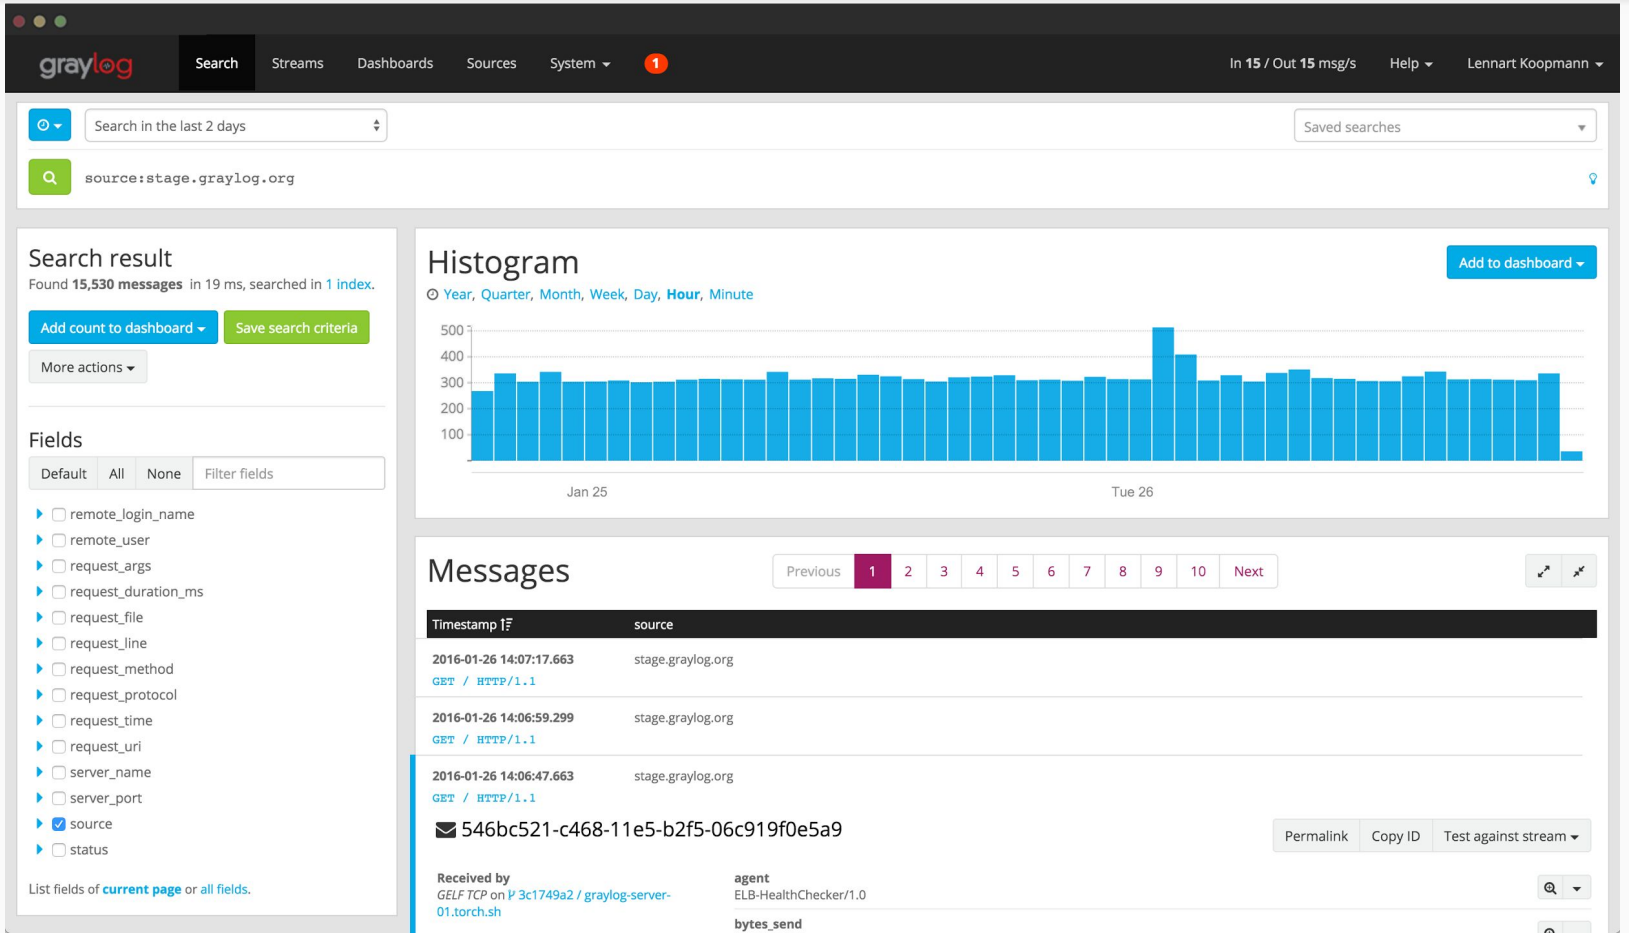

Dashboards

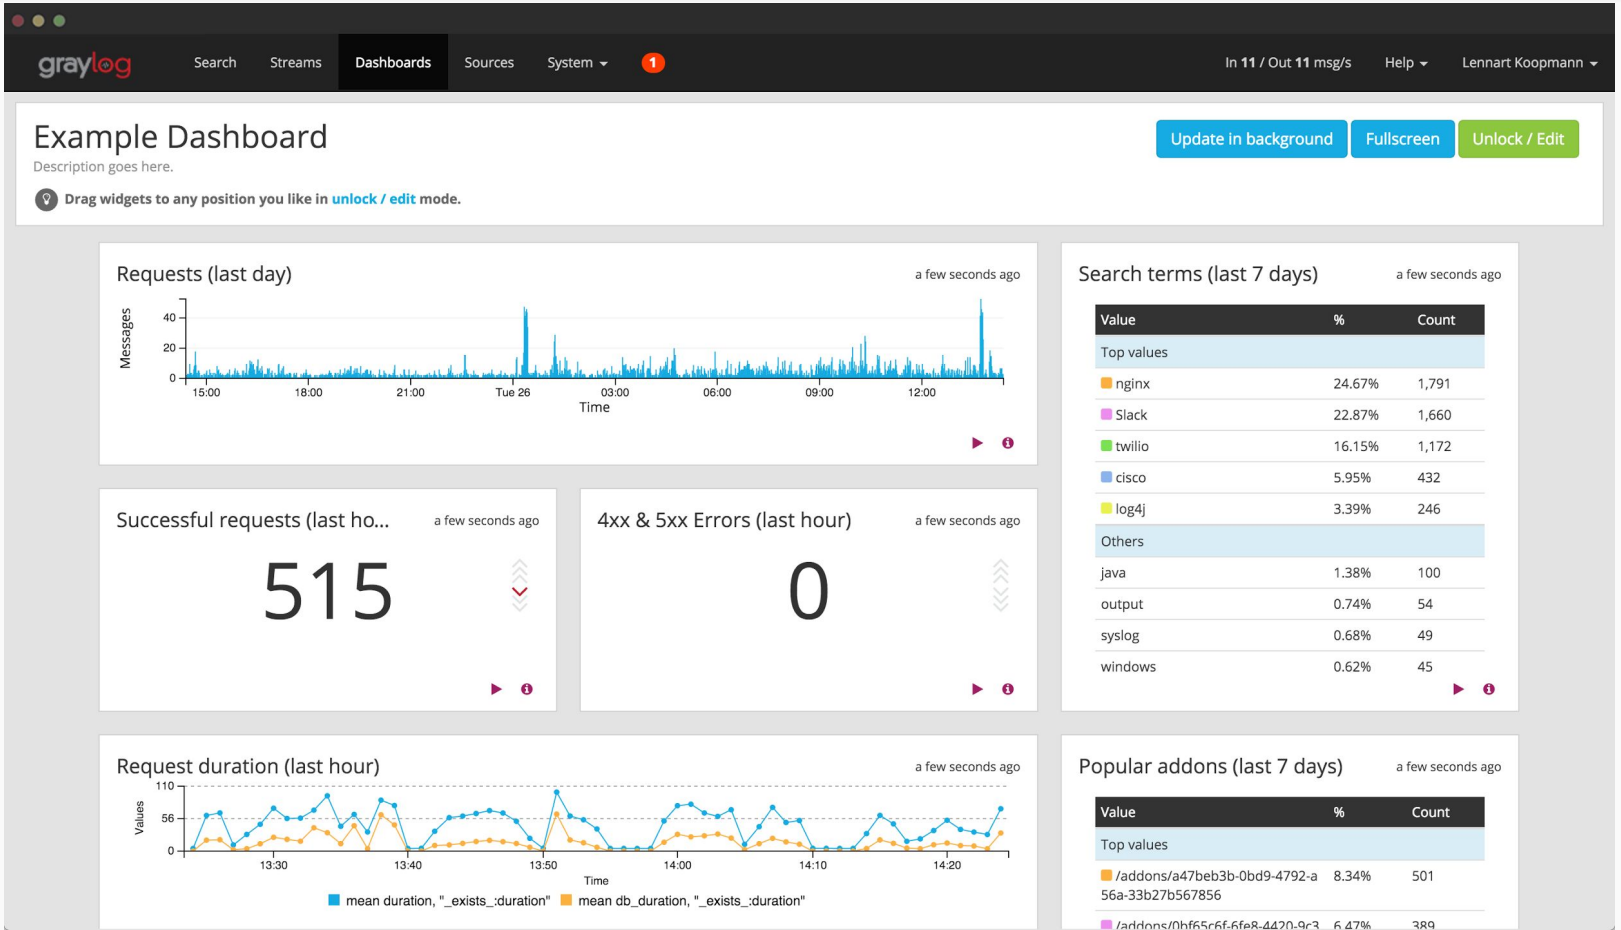
graylog
Search
Streams
Dashboards
Sources
System
1
In 11 / Out 11 msg/s
Help
Lennart Koopmann
Example Dashboard
Description goes here.
Drag widgets to any position you like in unlock / edit mode.
Update in background
Fullscreen
Unlock / Edit
Requests (last day)
a few seconds ago
Messages
40
20
0
15:00
18:00
21:00
Tue 26
03:00
06:00
09:00
12:00
Time
Search terms (last 7 days)
a few seconds ago
Value
%
Count
Top values
nginx 24.67% 1,791
Slack 22.87% 1,660
twilio 16.15% 1,172
cisco 5.95% 432
log4j 3.39% 246
Others
java 1.38% 100
output 0.74% 54
syslog 0.68% 49
windows 0.62% 45
Successful requests (last ho...
a few seconds ago
515
4xx & 5xx Errors (last hour)
a few seconds ago
0
Request duration (last hour)
a few seconds ago
Values
110
56
0
13:30
13:40
13:50
14:00
14:10
14:20
Time
mean duration, "_exists_:duration"
mean db_duration, "_exists_:duration"
Popular addons (last 7 days)
a few seconds ago
Value
%
Count
Top values
/addons/a47beb3b-0bd9-4792-a56a-33b27b567856 8.34% 501

Alerts

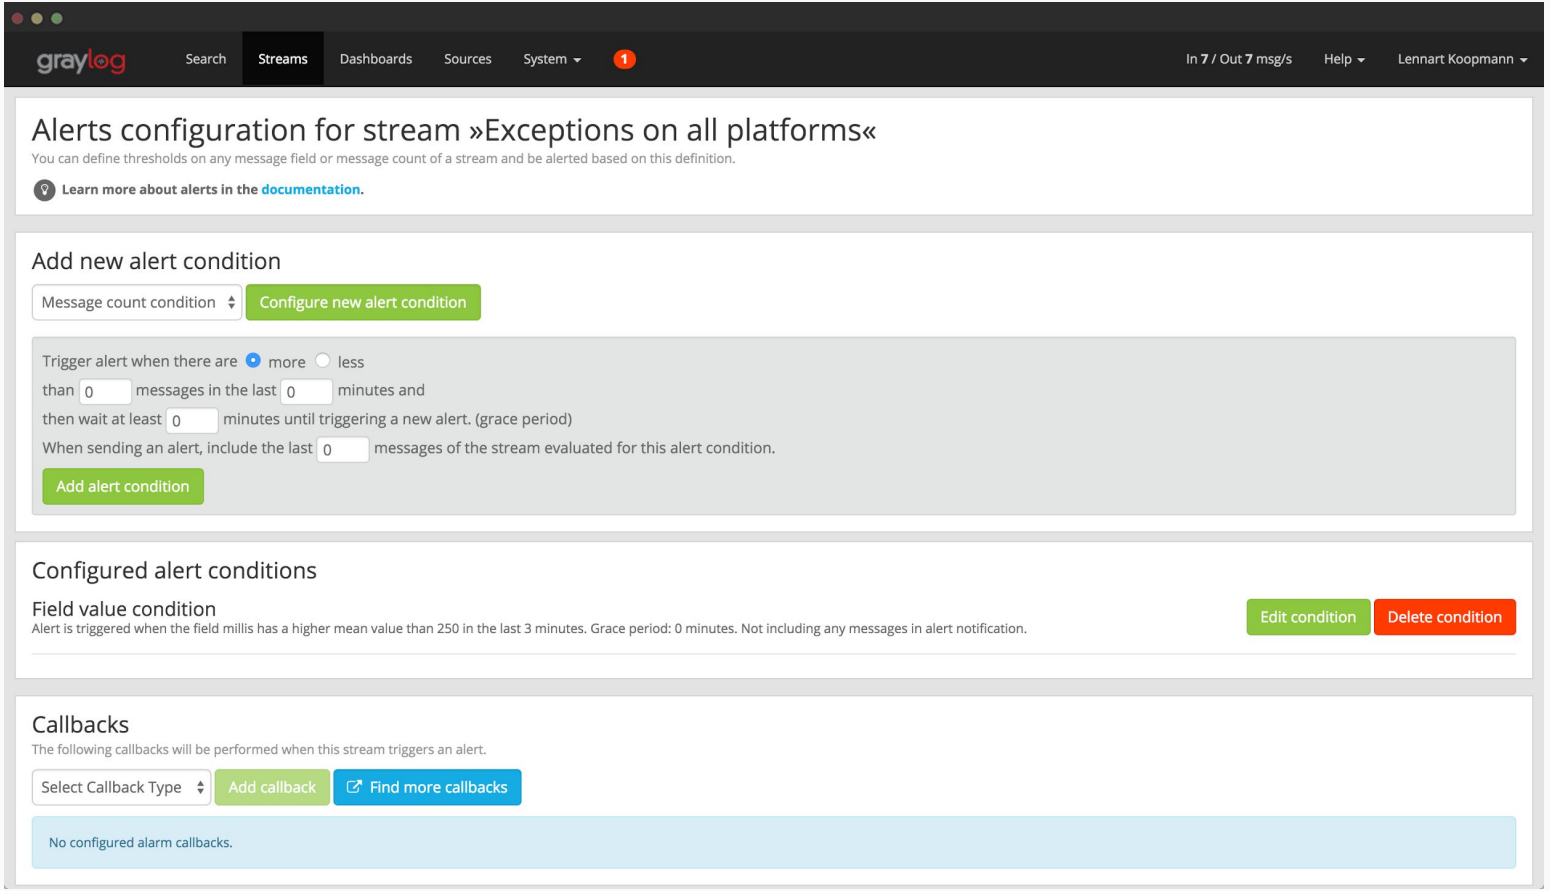
graylog
Search
Streams
Dashboards
Sources
System
1
In 7 / Out 7 msg/s
Help
Lennart Koopmann
Alerts configuration for stream »Exceptions on all platforms«
You can define thresholds on any message field or message count of a stream and be alerted based on this definition.
Learn more about alerts in the documentation.
Add new alert condition
Message count condition
Configure new alert condition
Trigger alert when there are more less
than 0 messages in the last 0 minutes and
then wait at least 0 minutes until triggering a new alert. (grace period)
When sending an alert, include the last 0 messages of the stream evaluated for this alert condition.
Add alert condition
Configured alert conditions
Field value condition
Alert is triggered when the field millis has a higher mean value than 250 in the last 3 minutes. Grace period: 0 minutes. Not including any messages in alert notification.
Edit condition
Delete condition
Callbacks
The following callbacks will be performed when this stream triggers an alert.
Select Callback Type
Add callback
Find more callbacks
No configured alarm callbacks.

API Browser

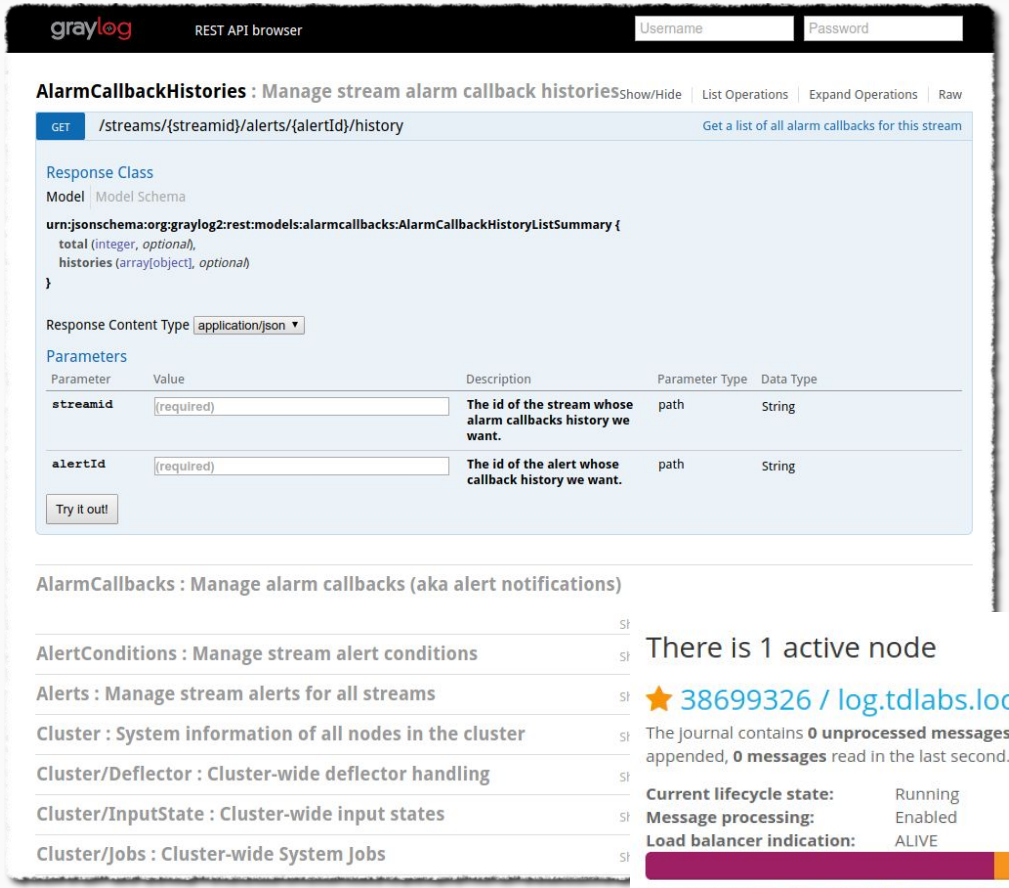
graylog
REST API browser
Username
Password
AlarmCallbackHistories : Manage stream alarm callback histories
Show/Hide
List Operations
Expand Operations
Raw
GET
/streams/{streamid}/alerts/{alertId}/history
Get a list of all alarm callbacks for this stream
Response Class
Model
Model Schema
urn:jsonschema:org:graylog2:rest:models:alarmcallbacks:AlarmCallbackHistoryListSummary {
total (integer, optional),
histories (array[object], optional)
}
Response Content Type
application/json
Parameters
Parameter
Value
Description
Parameter Type
Data Type
streamid
(required)
The id of the stream whose alarm callbacks history we want.
path
String
alertId
(required)
The id of the alert whose callback history we want.
path
String
Try it out!
AlarmCallbacks : Manage alarm callbacks (aka alert notifications)
AlertConditions : Manage stream alert conditions
Alerts : Manage stream alerts for all streams
Cluster : System information of all nodes in the cluster
Cluster/Deflector : Cluster-wide deflector handling
Cluster/InputState : Cluster-wide input states
Cluster/Jobs : Cluster-wide System Jobs

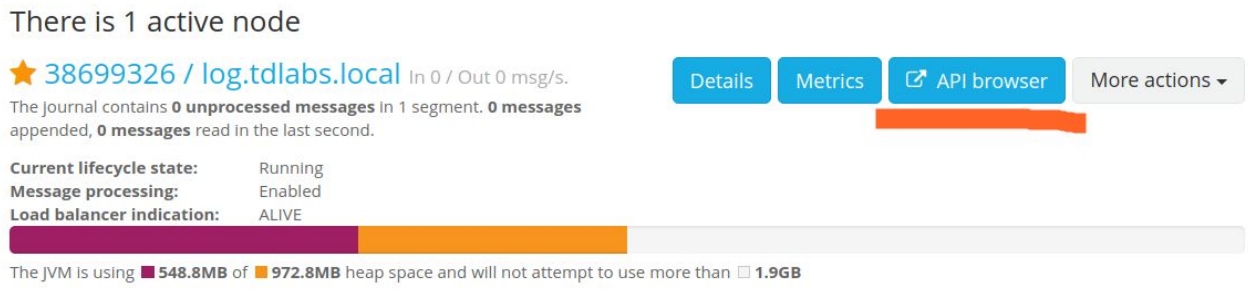
There is 1 active node
38699326 / log.tdlabs.local
In 0 / Out 0 msg/s.
Details
Metrics
API browser
More actions
The journal contains 0 unprocessed messages in 1 segment. 0 messages appended, 0 messages read in the last second.
Current lifecycle state: Running
Message processing: Enabled
Load balancer indication: ALIVE
The JVM is using 548.8MB of 972.8MB heap space and will not attempt to use more than 1.9GB

Outputs

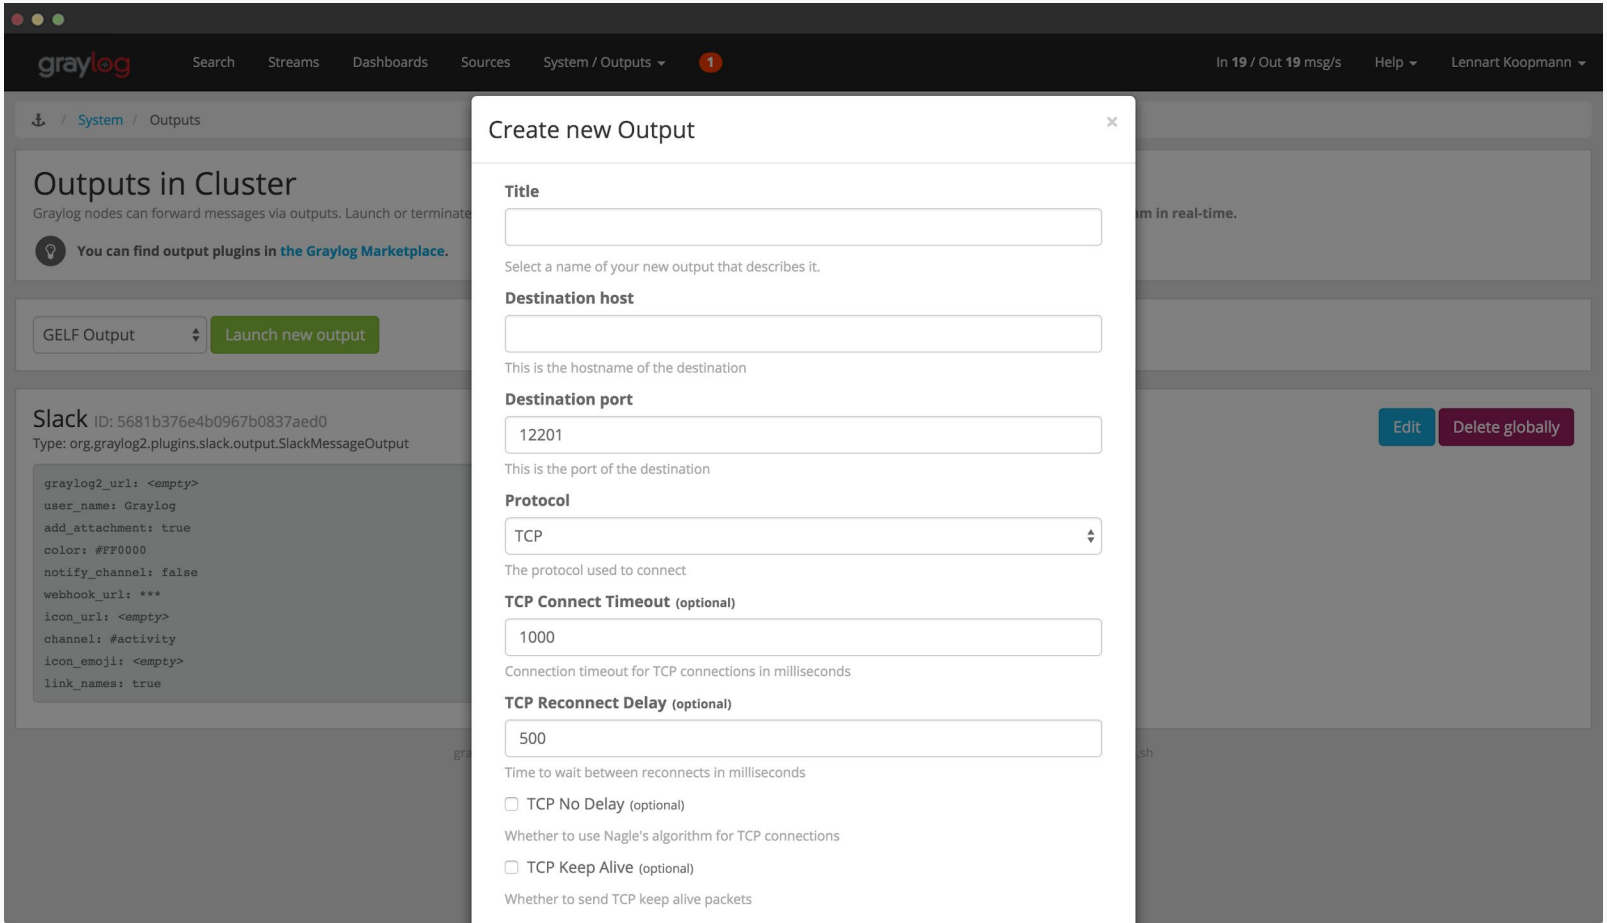
graylog
Search
Streams
Dashboards
Sources
System / Outputs
1
In 19 / Out 19 msg/s
Help
Lennart Koopmann
System / Outputs
Outputs in Cluster
Graylog nodes can forward messages via outputs. Launch or terminate
You can find output plugins in the Graylog Marketplace.
GELF Output
Launch new output
Slack ID: 5681b376e4b0967b0837aed0
Type: org.graylog2.plugins.slack.output.SlackMessageOutput
graylog2_url: <empty>
user_name: Graylog
add_attachment: true
color: #FF0000
notify_channel: false
webhook_url: ***
icon_url: <empty>
channel: #activity
icon_emoji: <empty>
link_names: true
Edit
Delete globally
Create new Output
Title
Select a name of your new output that describes it.
Destination host
This is the hostname of the destination
Destination port
12201
This is the port of the destination
Protocol
TCP
The protocol used to connect
TCP Connect Timeout (optional)
1000
Connection timeout for TCP connections in milliseconds
TCP Reconnect Delay (optional)
500
Time to wait between reconnects in milliseconds
TCP No Delay (optional)
Whether to use Nagle's algorithm for TCP connections
TCP Keep Alive (optional)
Whether to send TCP keep alive packets

Integrations

- Java
  - logstash-gelf Library
    - Support for multiple Logging Frameworks
    - Website http://logging.paluch.biz/
    - Github https://github.com/mp911de/logstash-gelf
    - Examples mp911de/logstash-gelf src/test/java/biz/paluch/logging/gelf
- .Net
  - gelf4net https://github.com/jjchiw/gelf4net
- Go
  - go-gelf https://github.com/Graylog2/go-gelf
- Windows
  - winlogbeat, Graylog Collector Sidecar
  - nxlog https://nxlog.co/products/nxlog-community-edition
- Linux
  - filebeat, Graylog Collector Sidecar
  - nxlog, syslog

DEMO

GELF & Java

Logback & GELF example configuration

```
<appender name="GELF" class="biz.paluch.logging.gelf.logback.GelfLogbackAppender">
    <host>${LOG_PROTO:-udp}:${LOG_HOSTNAME:-localhost}</host>
    <port>${LOG_PORT:-12201}</port>
    <version>1.1</version>
    <timestampPattern>yyyy-MM-dd HH:mm:ss,SSSS</timestampPattern>
    <maximumMessageSize>8192</maximumMessageSize>
    <facility>-</facility>
    <extractStackTrace>true</extractStackTrace>
    <filterStackTrace>true</filterStackTrace>
    <mdcProfiling>false</mdcProfiling>
    <additionalFields>org=tdlabs,ctx=demo,svc=hello-world-svc,env=test</additionalFields>
    <additionalFieldTypes>org=String,ctx=String,svc=String,env=String</additionalFieldTypes>
    <mdcFields>APP_STAGE</mdcFields>
    <dynamicMdcFields>svc_.*</dynamicMdcFields>
    <filter class="ch.qos.logback.classic.filter.ThresholdFilter">
        <level>${LOG_LEVEL_GELF:-INFO}</level>
    </filter>
</appender>

<root level="INFO">
    <appender-ref ref="GELF"/>
    <appender-ref ref="CONSOLE"/>
</root>
```

Log-Server Destination

StackTrace handling

Thread-Local **M**apped **D**iagnostic **C**ontext fields

logback.xml

Further reading

- Graylog 2.2 Design Documents
- Blog Post Monitoring Graylog
- Blog Post Processing 250GB Log Data / Day
- German Article in IT-Administrator 2015/09
- German Article in IT-Administrator 2015/10
- German Article in IT-Administrator 2015/11
- Youtube Windows Event log with Graylog

Questions


QUESTIONS
FOUND IN GOOGLE AUTOCOMPLETE
WHY DO WHALES JUMP
WHY ARE WITCHES GREEN
WHY ARE THERE MIRRORS ABOVE BEDS
WHY DO I SAY UH
WHY IS SEA SALT BETTER
WHY ARE THERE TREES IN THE MIDDLE OF FIELDS
WHY IS THERE NOT A POKEMON MMO
WHY IS THERE LAUGHING IN TV SHOWS
WHY ARE THERE DOORS ON THE FREEWAY
WHY ARE THERE SO MANY SHOWS ON E RUNNING
WHY AREN'T THERE ANY COUNTRIES IN ANTARCTICA
WHY ARE THERE SORRY SOUNDS IN MINECRAFT
WHY IS THERE KICKING IN MY STOMACH
WHY ARE THERE TWO SLASHES AFTER HTTP
WHY ARE THERE CELEBRITIES
WHY DO SNAKES EXIST
WHY DO OYSTERS HAVE PEARLS
WHY ARE DUCKS CALLED DUCKS
WHY DO THEY CALL IT THE CLAP
WHY ARE KYLE AND CARTMAN FRIENDS
WHY IS THERE AN ARROW ON FANG'S HEAD
WHY ARE TEXT MESSAGES BLUE
WHY ARE THERE MUSTACHES ON CLOTHES
WHY ARE THERE MUSTACHES ON CARS
WHY ARE THERE MUSTACHES EVERYWHERE
WHY ARE THERE SO MANY BIRDS IN OHIO
WHY IS THERE SO MUCH RAIN IN OHIO
WHY IS OHIO WEATHER SO WEIRD
WHY ARE THERE MALE AND FEMALE BIKES
WHY ARE THERE SQUIRRELS
WHY ARE THERE TINY SPIDERS IN MY HOUSE
WHY DO SPIDERS COME INSIDE
WHY ARE THERE HUGE SPIDERS IN MY HOUSE
WHY ARE THERE LOTS OF SPIDERS IN MY HOUSE
WHY ARE THERE SPIDERS IN MY ROOM
WHY ARE THERE SO MANY SPIDERS IN MY ROOM
WHY DO SPIDER BITES ITCH
WHY IS DYING SO SCARY
WHY DO KNEES CLICK
WHY AREN'T THERE E-GARDENS
WHY IS ISOLATION BAD
WHY DO BOYS LIKE ME
WHY DON'T BOYS LIKE ME
WHY IS LYING GOOD
WHY DO TREES DIE
WHY DO TESTICLES MOVE
WHY ARE THERE PSYCHICS
WHY ARE HATS SO EXPENSIVE
WHY DO YOUR BOOGERS HURT
WHY ARE THERE SLAVES IN THE BIBLE
WHY DO TWINS HAVE DIFFERENT FINGERPRINTS
WHY ARE AMERICANS AFRAID OF DRAGONS
WHY IS HTTPS CROSSED OUT IN RED
WHY IS THERE A LINE THROUGH HTTPS
WHY IS HTTPS IMPORTANT
WHY AREN'T MY ARMS GROWING
WHY AREN'T ECONOMISTS RICH
WHY DO AMERICANS CALL IT SOCCER
WHY ARE MY EARS RINGING
WHY ARE THERE SO MANY AVENGERS
WHY ARE THE AVENGERS FIGHTING THE X-MEN
WHY IS WOLVERINE NOT IN THE AVENGERS
WHY ARE THERE SO MANY CROWS IN ROCHESTER, MN
WHY IS PSYCHIC WEAK TO BUG
WHY DO CHILDREN GET CANCER
WHY IS POSEIDON ANGRY WITH ODYSSEUS
WHY IS THERE ICE IN SPACE
WHY ARE THERE ANTS IN MY LAPTOP
WHY IS EARTH TILTED
WHY IS SPACE BLACK
WHY IS OUTER SPACE SO COLD
WHY ARE THERE PYRAMIDS ON THE MOON
WHY IS NASA SHUTTING DOWN
WHY ARE THERE GHOSTS
WHY IS THERE AN OWL IN MY BACKYARD
WHY IS THERE AN OWL OUTSIDE MY WINDOW
WHY IS THERE AN OWL ON THE DOLLAR BILL
WHY DO OWLS ATTACK PEOPLE
WHY ARE AK 47S SO EXPENSIVE
WHY ARE THERE HELICOPTERS CIRCLING MY HOUSE
WHY ARE THERE GODS
WHY ARE THERE TWO SPOOKS
WHY ARE MY BOOBS ITCHY
WHY ARE CIGARETTES LEGAL
WHY ARE THERE DUCKS IN MY POOL
WHY IS JESUS WHITE
WHY IS THERE LIQUID IN MY EAR
WHY DO Q-TIPS FEEL GOOD
WHY DO GOOD PEOPLE DIE
WHY ARE DOGS AFRAID OF FIREWORKS
WHY IS THERE NO KING IN ENGLAND
WHY IS MT VESUVIUS THERE
WHY DO THEY SAY T MINUS
WHY ARE THERE OBELISKS
WHY ARE WRESTLERS ALWAYS WET
WHY ARE OCEANS BECOMING MORE ACIDIC
WHY IS ARWEN DYING
WHY AREN'T MY QUAIL LAYING EGGS
WHY AREN'T MY QUAIL EGGS HATCHING
WHY AREN'T THERE ANY FOREIGN MILITARY BASES IN AMERICA
WHY IS LIFE SO BORING
WHY IS SEX SO IMPORTANT
WHY AREN'T THERE GUNS IN HARRY POTTER
WHY IS STEALING WRONG
WHY ARE THERE FEMALE MR MIMES
WHY IS THERE HELL IF GOD FORGIVES
WHY IS GPS FREE

Best Practices: Log Message Enrichment

- System Context
  - Where did the log message originate?
  - → Associate context information with the log Message
- Request Context
  - Who processed the message?
  - Follow the request processing through multiple layers (Request Id)
  - ... or even accros multiple nodes (Trace Id) → http://zipkin.io
- Audit Information
  - Which user did produce the log message?
  - Beware of privacy law!
- First Failure Data Capture
  - Create a unique id for each particular error instance
  - → Makes it easier to refer to the error

System Context

- **`source`** — *Host* — dborac1a.db.internal.acme.com
- `org` — *Organization / Tenant* — acme, customer1, tdlabs
- `ctx` — *Context / System Boundary* — idm, net, accounting, clearing
- `env` — *Environment* — dev, local, test, qa, prod
- **`svc`** — *Logical Service name* — sso, booking, sla-monitoring
- `inst` — *Service Instance* — 1, 1a, 2b

Best Practices: Log Message Enrichment

Request Context

- **`rid`** *Request Id* 6caae423-64f8-326d
- `tid` *Trace Id* 12321-23231-2133-23

Audit Information

- **`usr_id`** *Global/Tenant User Id* c8609423-66d8-485d
- `usr_name` *Tenant User Name* ameier

First Failure Data Capture

- **`err_id`** *UUID per Error* aa2a-4e10609b95a1
- **`err_code`** *Logical Error Code* BILLING_ERROR_BANK_IFACE_UNAVIL

Thomas Darimont
Software Architect >eurodata AG

Java User Group
Saarland

t.darimont@eurodata.de
@thomasdarimont